

---

# Reducing Deep Learning training times with parallelism strategies and TensorFlow Distribute.

Juan Manuel Muñoz Betancur



# PRESENTATION OUTLINE

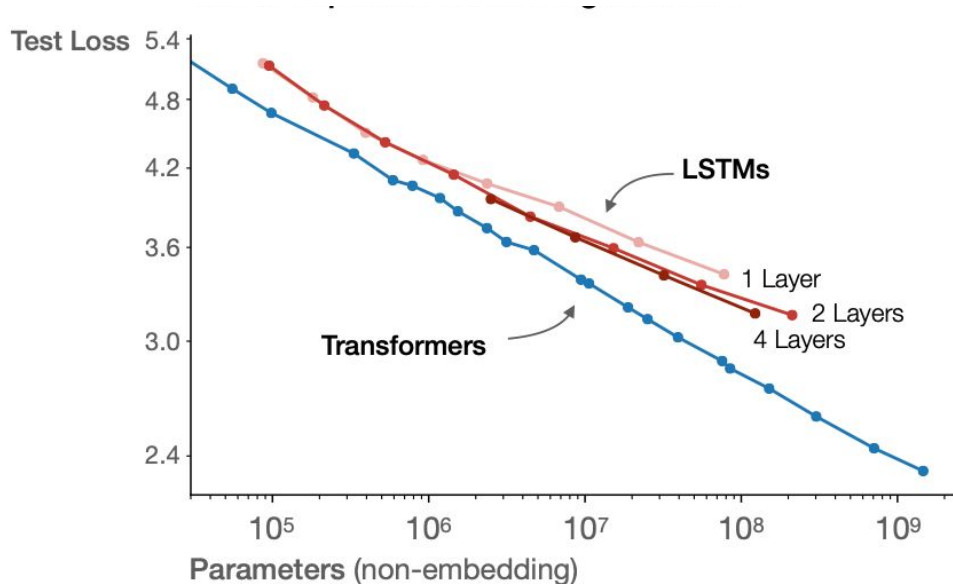
- Introduction/Motivation.
- Common types of parallelism.
  - Model Parallelism
  - Data Parallelism
- Tensorflow Demo.



# Motivation

Get results faster and iterate quickly in your small/medium/large projects without additional costs.

State of the art models: GPT-3 has 175 Billion parameters, to train it even in the best SINGLE GPU available in the market would take about 355 years.



# Motivation

To train a model you need to take into account the size of the parameters, the forward pass activations (saving them for the backprop) and the gradients of the weights.

## Params size

InceptionV3 24M params  
@ fp64 =  $24 \times 10^6 \times 8$  bytes  
**= 192 MB**

## Activations size

InceptionV3 @ fp64 batch\_size=256  
~ = 256 (4 x  $10^6$  x 8 bytes )  
**= 32MB \* 256 = 8GB**

## Params size

GPT-3 175B parameters @fp16  
=  $175 \times 10^9 \times 2$  bytes  
**= 316.2 GB**

## Activations size

GPT-3 ~TBs



# COMMON TYPES OF PARALLELISM


- Model Parallelism.
  - Pipeline Parallelism
- Data Parallelism.




# Model Parallelism (Large model training)

**split the model across devices**

each device runs a fragment of the model




## Pipeline Parallelism



## Distributed Tensor Computation


GPU 1



GPU 2




# Model Parallelism: Naive







## Model Parallelism: Naive



# Model Parallelism: Naive




# Pipeline Parallelism




**key idea:** split mini-batch into sequential micro-batches

# Pipeline Parallelism




Micro-Batch 1


# Pipeline Parallelism




# Pipeline Parallelism



# Pipeline Parallelism



# Pipeline Parallelism





# Pros and cons of model parallelism

## Pros:

- Can train bigger models.
- Implemented on Pytorch.

## Cons:

- Not found in the distribution strategy of default libraries such as Tensorflow. (Mesh Tensorflow)
- Tricky to design an implement.




# Data Parallelism (Large Batch Training)

**split the data across devices**


each device sees a fraction of the batch

each device replicates the model


each device replicates the optimizer



# Data Parallelism




**GPU 1**




**GPU 2**

GPUs could be on same or multiple nodes

## Get a batch of data




**GPU 1**




**GPU 2**




## Split batch across devices




**GPU 1**




**GPU 2**




# Parallel forward passes




**GPU 1**




**GPU 2**




# Parallel forward passes




**GPU 1**




**GPU 2**



## Backpropagate gradients




**GPU 1**



**GPU 2**




# Backpropagate gradients



**GPU 1**


Gradients GPU1




**GPU 2**

Gradients GPU2


# Share gradients among GPUs and update



Start the next step with a new  
minibatch of data




GPU 1



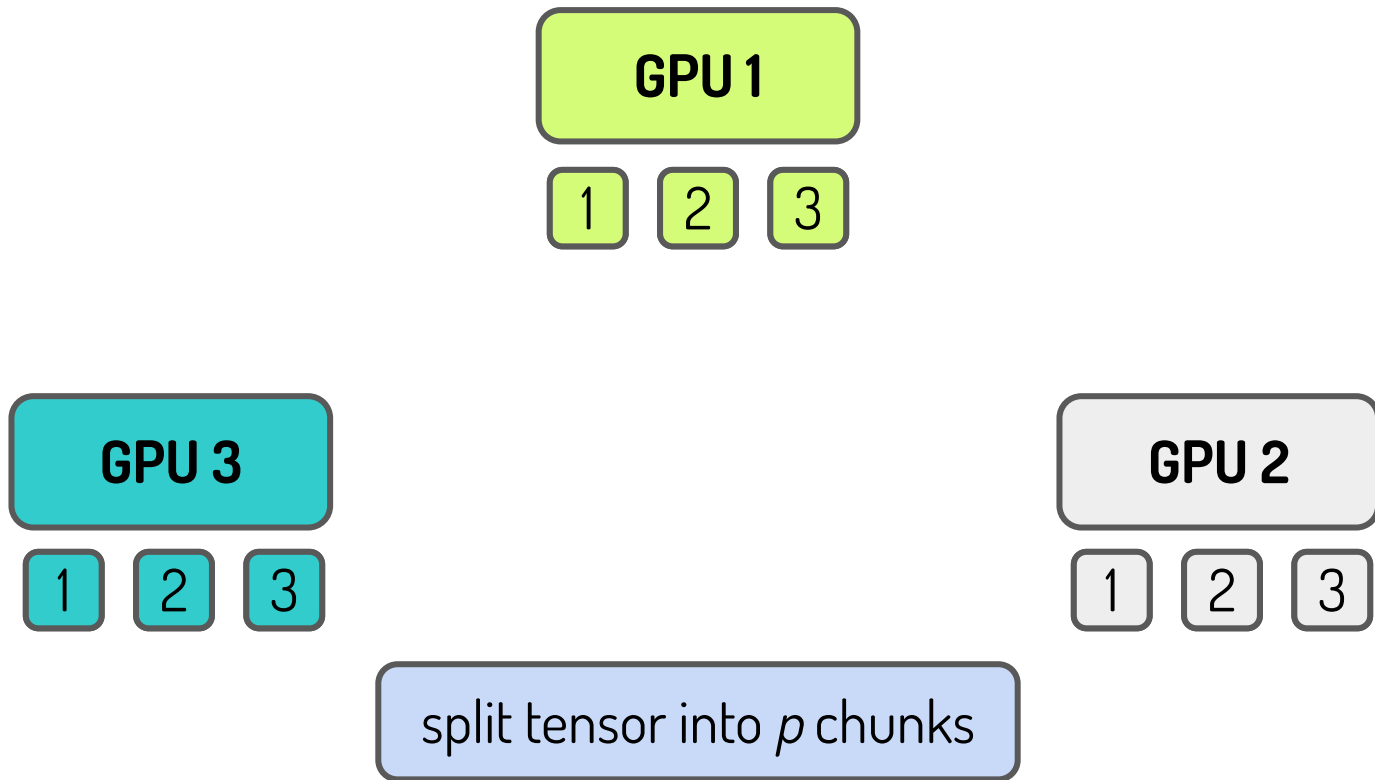
GPU 2

**all parameters stay synchronized!**


# So, what's All-Reduce?




# Ring All-Reduce




# Ring All-Reduce



# Ring All-Reduce




# Ring All-Reduce

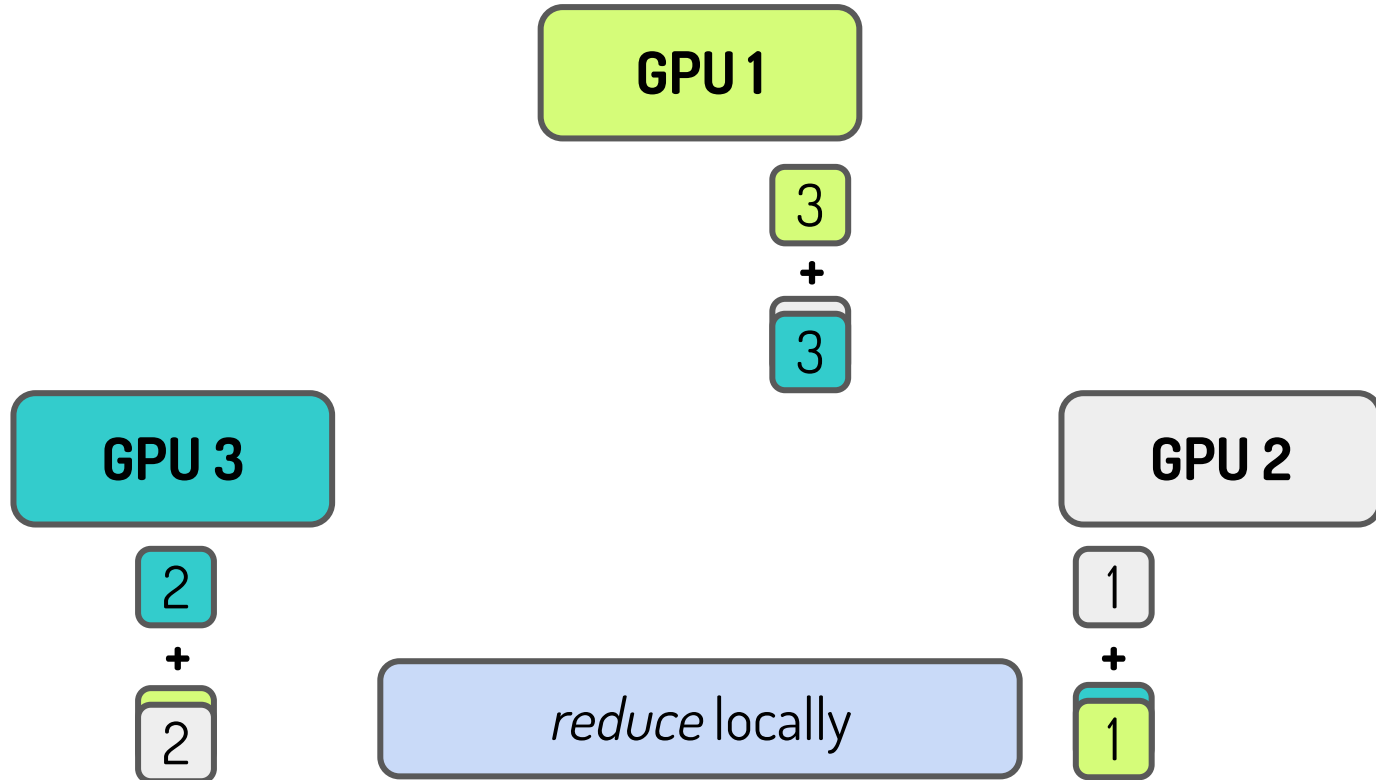





# Ring All-Reduce



# Ring All-Reduce




# Ring All-Reduce




everyone has a chunk of the result after  $(p-1)$  share-reduce steps


# Ring All-Reduce




# Ring All-Reduce



# Ring All-Reduce



# Ring All-Reduce



# Ring All-Reduce Advantages

Naive method =  $p$  senders  $\times (p - 1)$  receivers  $\times o(n)$  tensor =  **$o(np^2)$**   
everyone does  **$o(np)$**  work.

Manager node method =  $(p-1) \times 2$  transfers  $\times o(n)$  tensor =  **$o(np)$**   
manager does  **$o(np)$**  work

Ring All-Reduce =  $p$  senders  $\times 1$  receiver  $\times o(n/p)$  tensor  $\times (p-1)$  rounds  $\times 2$  phases =  **$o(np)$**   
everyone does equal  **$o(n)$**  work (independent of  $p$ )



# Before the demo, what strategies are implemented in Tensorflow

| Training API         | MirroredStrategy | TPUStrategy   | MultiWorkerMirroredStrategy | CentralStorageStrategy | ParameterServerStrategy    |
|----------------------|------------------|---------------|-----------------------------|------------------------|----------------------------|
| Keras API            | Supported        | Supported     | Supported                   | Experimental support   | Supported planned post 2.4 |
| Custom training loop | Supported        | Supported     | Supported                   | Experimental support   | Experimental support       |
| Estimator API        | Limited Support  | Not supported | Limited Support             | Limited Support        | Limited Support            |



# Demos

Using data parallelism with  
`model.fit`

Using data parallelism with  
a custom training loop.

(<https://github.com/juanma9613/Reducing-deep-learning-training-times-Pycon2021>)



# THANK YOU!

- <https://www.linkedin.com/in/juan-manuel-munoz-betancur/>
- [jmunozb@eafit.edu.co](mailto:jmunozb@eafit.edu.co)

Link to demo:

- <https://github.com/juanma9613/Reducing-deep-learning-training-times-Python2021>

## Factored is hiring!